



A Survey on Deep Reinforcement Learning

PhD Qualifying Examination


Siyi LI

Supervisor:

Prof. Dit-Yan Yeung

Background

- Deep learning methods have making major advances in solving many low-level perceptual tasks.



See (visual object recognition)



Read (text understanding)



Hear (speech recognition)

Background

- More sophisticated tasks that involve decision and planning require an higher level of intelligence.
- Real Artificial Intelligence system also requires the ability of reasoning, thinking and planning.



Playing Atari game




Robotic navigation

Limitations of Deep Learning

- Supervised learning assumptions
 - Training and testing instances are **i.i.d** variables
 - Training data are **labeled** data with **strong supervision**
- Reality of most real-world tasks
 - Strong supervision is **expensive** and **scarce**
 - **Sequential interactive** process violates the i.i.d assumption

Reinforcement Learning (RL) in a nutshell

- RL is a general-purpose framework for decision making
 - RL is for an **agent** to **act** with an **environment**
 - Each **action** influences the agent's future **state**
 - Feedback is given by a scalar **reward** signal
 - Goal: **select actions to maximise future reward**



Deep RL: Deep Learning + RL


- Traditional RL **action** approaches have been limited to domains with **low-dimensional state** spaces or **handcrafted features**
- By combining deep learning and RL, we want to embrace both the representation power of deep learning and generalization ability from RL
 - RL defines the **objective**
 - Deep Learning learns the **representation**

Outline

- Introduction to Deep Learning
- Introduction to Reinforcement Learning (RL)
- Value-Based Deep RL
- Policy-Based Deep RL
- Other Deep RL Extensions
- Deep RL Applications

Deep Representations

- A **deep representation** is a composition of many functions, where each composition level is learning representations at different level of abstraction



- The weights are learned using **backpropagation** by chain rule

The diagram illustrates the backpropagation process through the hidden states. It shows the flow of gradients from the loss l back to the weights w_1, \dots, w_n . The top row shows the gradient flow from l to y and then through the hidden states h_n, \dots, h_1 to x . The bottom row shows the gradient flow from h_n and h_1 down to the weights w_n and w_1 respectively. The middle row shows the gradient flow from h_1 to w_1 and from h_n to w_n .

$$\begin{array}{ccccccc} \frac{\partial l}{\partial x} & \xleftarrow{\frac{\partial h_1}{\partial x}} & \frac{\partial l}{\partial h_1} & \xleftarrow{\frac{\partial h_2}{\partial h_1}} & \dots & \xleftarrow{\frac{\partial h_n}{\partial h_{n-1}}} & \frac{\partial l}{\partial h_n} \xleftarrow{\frac{\partial y}{\partial h_n}} \frac{\partial l}{\partial y} \\ & & \downarrow \frac{\partial h_1}{\partial w_1} & & & & \downarrow \frac{\partial h_n}{\partial w_n} \\ & & \frac{\partial l}{\partial w_1} & & \dots & & \frac{\partial l}{\partial w_n} \end{array}$$

Deep Neural Network

- A deep neural network typically consists of:

- Linear transformations

$$h_{k+1} = Wh_k$$

- Nonlinear activation functions

$$h_{k+1} = \sigma(h_k)$$

$$\sigma(\cdot) = \tanh(\cdot), \frac{1}{1 + \exp(\cdot)}, \dots$$

- Loss function on the output



- Mean squared error: $I = ||y - y^*||^2$
- Log likelihood: $I = \log P(y^*)$

Training by Stochastic Gradient Descent

- Sample gradient of expected loss $L(\mathbf{w}) = \mathbb{E}[l]$ (better efficiency for large data)

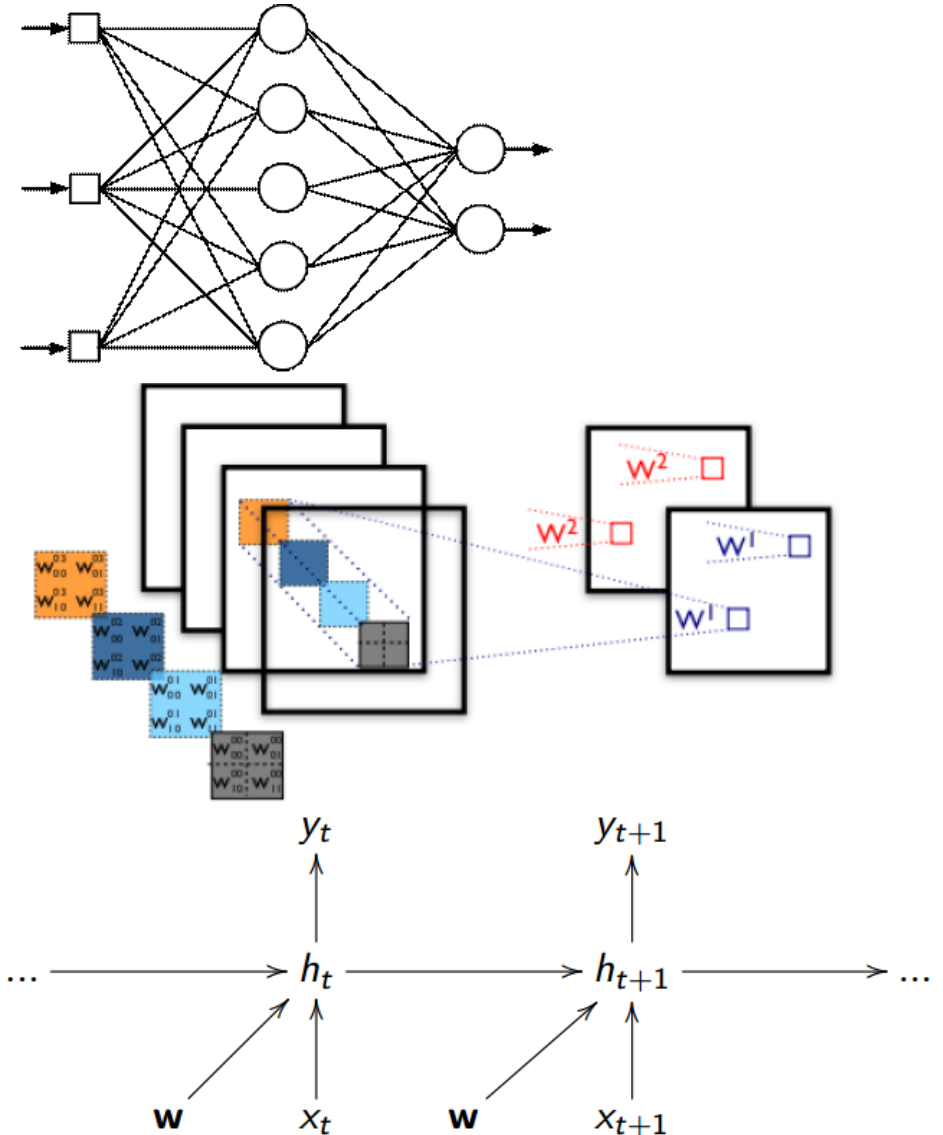
$$\frac{\partial l}{\partial \mathbf{w}} \sim \mathbb{E} \left[\frac{\partial l}{\partial \mathbf{w}} \right] = \frac{\partial L(\mathbf{w})}{\partial \mathbf{w}}$$

- Adjust \mathbf{w} down the sampled gradient



Deep Learning Models

- Multilayer perceptrons (MLPs)
 - Fully-connected
- Convolutional neural networks (CNNs)
 - Weight sharing between local regions
- Recurrent neural networks (RNNs)
 - Weight sharing between time-steps



Outline

- Introduction to Deep Learning
- **Introduction to Reinforcement Learning (RL)**
- Value-Based Deep RL
- Policy-Based Deep RL
- Other Deep RL Extensions
- Deep RL Applications

Markov Decision Processes (MDPs)

- MDPs formally describe an environment for RL, where the environment is fully observable

Definition


An MDP is a tuple $(\mathcal{S}, \mathcal{A}, f, R)$ consisting of:

- \mathcal{S} : The **state** space. In MDPs, the **state** is a sufficient statistic of the future.
- \mathcal{A} : The **action** space.
- $f: \mathcal{S} \times \mathcal{A} \times \mathcal{S} \mapsto [0, \infty)$: The **state transition probability** density function.

$$P(s_{k+1} \in \mathcal{S}_{k+1} | s_k, a_k) = \int_{\mathcal{S}_{k+1}} f(s_k, a_k, s') ds'$$

- $R: \mathcal{S} \times \mathcal{A} \times \mathcal{S} \mapsto \mathbb{R}$: The **reward** function.

$$r_k = R(s_k, a_k, s_{k+1})$$



Policy

A **policy** is the behavior of the agent.

- Stochastic policy $\pi: \mathcal{S} \times \mathcal{A} \mapsto [0, \infty)$.

$$P(a|s) = \pi(s, a)$$

- Deterministic policy $\pi: \mathcal{S} \mapsto \mathcal{A}$.

$$a = \pi(s)$$

Expected Return

- The goal of RL is to find the policy which maximizes the **expected return**

$$J(\pi) = \mathbb{E}\{g(r_0, r_1, \dots) | \pi\}.$$


- In most cases, the function **g** is either the **discounted sum of rewards** or the **average reward**

- Discounted reward

$$g(r_0, r_1, \dots) = \sum_{k=0}^{\infty} \gamma^k r_k$$

- Average reward

$$g(r_0, r_1, \dots) = \lim_{n \rightarrow \infty} \frac{1}{n} \sum_{k=0}^{n-1} r_k$$



Value Function

- Consider the **discounted reward** case

$$\begin{aligned} J(\pi) &= \mathbb{E} \left\{ \sum_{k=0}^{\infty} \gamma^k r_k \middle| d_0, \pi \right\} \\ &= \int_S d^\pi(s) \int_{\mathcal{A}} \pi(s, a) \int_S f(s, a, s') R(s, a, s') ds' da ds. \end{aligned}$$

$$d^\pi(s) = \sum_{k=0}^{\infty} \gamma^k p(s_k = s | d_0, \pi)$$

- A **value function** is the **prediction** of the above expected return
- Two definitions exist for the value function

- State value function**

$$V^\pi(s) = \mathbb{E} \left\{ \sum_{k=0}^{\infty} \gamma^k r_k \middle| s_0 = s, \pi \right\}$$

- State-action value function**

$$Q^\pi(s, a) = \mathbb{E} \left\{ \sum_{k=0}^{\infty} \gamma^k r_k \middle| s_0 = s, a_0 = a, \pi \right\}$$



$$V^\pi(s) = \mathbb{E} \{ Q^\pi(s, a) | a \sim \pi(s, \cdot) \}$$

Bellman Equation and Optimality

- Value functions decompose into **Bellman equations**, i.e., the value functions can be decomposed into immediate reward plus discounted value of successor state

$$V^\pi(s) = \mathbb{E} \{ R(s, a, s') + \gamma V^\pi(s') \}$$

$$Q^\pi(s, a) = \mathbb{E} \{ R(s, a, s') + \gamma Q^\pi(s', a') \}$$

- An **optimal** value function is the maximum achievable value.

$$V^*(s) = \max_{\pi} V^\pi(s) \quad Q^*(s, a) = \max_{\pi} Q^\pi(s, a)$$

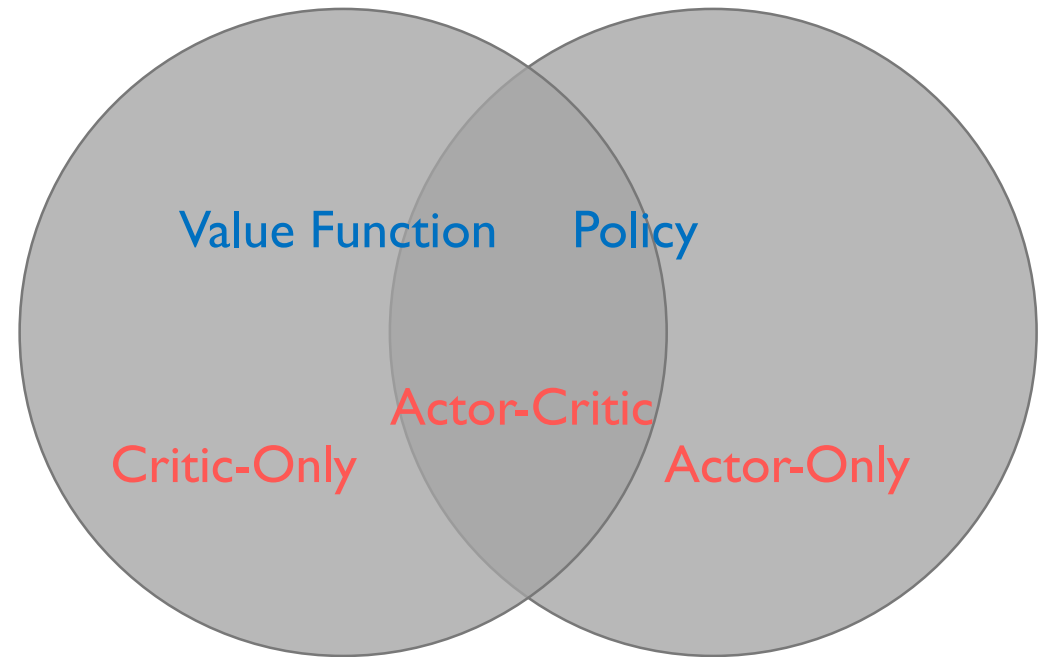
- Optimality for value functions are governed by the **Bellman optimality equations**.

$$V^*(s) = \max_a \mathbb{E} \{ R(s, a, s') + \gamma V^*(s') \}$$

$$Q^*(s, a) = \mathbb{E} \left\{ R(s, a, s') + \max_{a'} \gamma Q^*(s', a') \right\}$$

Approaches to RL

- **Critic-Only** Methods
 - Learnt value function
 - Implicit policy
- **Actor-Only** Methods
 - No value function
 - Learnt policy
- **Actor-Critic** Methods
 - Learnt value function
 - Learnt policy



Actor and *critic* are synonyms for the *policy* and *value* function.

Critic-Only

- A value function defines optimal policy.
- In critic-only methods, policy can be derived by selecting *greedy actions*

$$\pi^*(s) = \arg \max_a \mathbb{E} \{ R(s, a, s') + \gamma V^*(s') \}$$

$$\pi^*(s) = \arg \max_a Q^*(s, a).$$


- Finding optimal policy from state-action value function is direct.
- Finding optimal policy from state value function is more complicated (requires the knowledge of *transition model*)

Critic-Only

- Dynamic programming-based methods
 - Policy iteration
 - Value iteration
- Monte Carlo (MC) methods
- Temporal difference (TD) learning methods
 - $TD(\lambda)$
 - Q-learning
 - SARSA

Dynamic Programming (DP)

- DP methods require a model of the **state transition density** function f and the **reward function** R to calculate the state value function
- DP is *model-based*
- **Policy evaluation**: updates the value function for the current policy
- **Policy improvement**: improve the policy by acting according to the current value function
- Typical methods:
 - Policy iteration
 - Alternates between policy evaluation and policy improvement
 - Value iteration
 - No explicit policy
 - Directly update value function



Monte Carlo Methods


- MC methods learn directly from **episodes** of experience
- MC is **model-free**: no knowledge of transitions/rewards
- MC learns from **complete** episodes.
 - All episodes must terminate
- MC uses simple idea: **empirical mean return**
- The estimates are **unbiased** but have **high variance**.

Temporal Difference (TD) Learning

- TD methods learn directly from **episodes** of experience
- TD is **model-free**: no knowledge of transitions/rewards
- TD learns from **incomplete** episodes, it can learn online after every step
- TD uses temporal errors to update value function

$$V'(s) = V(s) + \alpha (R(s, a, s') + \gamma V(s') - V(s))$$

$$Q'(s, a) = Q(s, a) + \alpha (R(s, a, s') + \gamma Q(s', a') - Q(s, a))$$



SARSA

$$Q'(s, a) = Q(s, a) + \alpha \left(R(s, a, s') + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)$$



Q-learning

- The estimates are **biased** but have **low variance**.

Actor-Only

- Critic-only methods do not scale well to high-dimensional or continuous action spaces, since selecting greedy actions is **computationally intensive**.
- Actor-only methods work with a parameterized family of policies over which optimization procedures can be used directly.
- Advantages
 - Better convergence properties
 - Effective in high-dimensional or continuous action spaces
 - Can learn stochastic policies
- Disadvantages
 - Evaluating a policy is inefficient and of high variance

Policy Gradient

- Given policy π_θ with parameters θ , the goal is find best θ to maximize the expected return
- Can use gradient descent

$$\theta_{k+1} = \theta_k + \alpha_{a,k} \nabla_\theta J(\theta_k)$$

- Policy gradient

$$\nabla_\theta J(\theta) = \frac{\partial J}{\partial \pi_\theta} \frac{\partial \pi_\theta}{\partial \theta} \leftarrow$$

- How to **estimate** the **policy gradient**?
 - Finite-difference methods
 - Likelihood ratio methods

Finite-Difference Methods

- Idea is simple, i.e., to vary the policy parameterization by small increments and evaluate the cost by rollouts
- Estimate k-th partial derivative of object function

$$\frac{\partial J(\theta)}{\partial \theta_k} \approx \frac{J(\theta + \epsilon u_k) - J(\theta)}{\epsilon}$$

- Simple, works even if policy is not differentiable
- Noisy, inefficient

Likelihood Ratio Methods (REINFORCE)

- We can formulate the expected return from the view of **trajectories** generated by rollouts

$$\tau \sim p(\tau|\theta) \quad J^\tau = \sum_{k=0}^H \gamma^k r_k \quad J(\theta) = \int_{\mathbb{T}} p(\tau|\theta) J^\tau d\tau$$

- Use **likelihood ratios** to compute the policy gradient

$$\nabla_{\theta} J(\theta) = \int_{\mathbb{T}} \nabla_{\theta} p(\tau|\theta) J^\tau d\tau$$

$$= \int_{\mathbb{T}} p(\tau|\theta) \nabla_{\theta} \log p(\tau|\theta) J^\tau d\tau$$

$$= \mathbb{E} \left\{ \underbrace{\nabla_{\theta} \log p(\tau|\theta)}_{\text{likelihood ratios}} J^\tau \right\} . \quad \leftarrow \nabla_{\theta} \log p(\tau|\theta) = \sum_{k=0}^H \nabla_{\theta} \log \pi_{\theta}(s_k, a_k)$$

Do not need to compute the system dynamics

$$\nabla_{\theta} J(\theta) = \mathbb{E} \left\{ \left(\sum_{k=0}^H \nabla_{\theta} \log \pi_{\theta}(s_k, a_k) \right) J^\tau \right\}$$

Likelihood Ratio Methods (REINFORCE)

- The above computation of policy gradient can be further reduced by replacing the **trajectory return** by the **state-action value** function

$$\nabla_{\theta} J(\theta) = \mathbb{E} \left\{ \sum_{k=0}^H \nabla \log \pi_{\theta}(s_k, a_k) Q^{\pi}(s_k, a_k) \right\}$$

- The trajectory return or the state-action value can be estimated by the return v_t obtained from Monte Carlo rollouts
- Thus the estimated policy gradient may have **large variance**
- In practice, subtracting a **baseline** from the trajectory return or the state-action value helps a lot

Actor-Critic

- Critic-only
 - Pros: low variance
 - Cons: difficult for continuous action domains
- Actor-only
 - Pros: easy to handle continuous actions
 - Cons: high variance
- Actor-critic combines the advantages of actor-only and critic-only methods
 - Actions are generated by the parameterized actor
 - The critic supplies the actor with low variance gradient estimates

Policy Gradient Theorem

- Actor-critic methods rely on the following **policy gradient theorem**

Theorem

(Policy Gradient): For any MDP, in either the discounted reward or average reward setting, the policy gradient is given by

$$\nabla_{\theta} J(\theta) = \int_{\mathcal{S}} d^{\pi}(s) \int_{\mathcal{A}} \nabla_{\theta} \pi_{\theta}(s, a) Q^{\pi}(s, a) da ds$$

with $d^{\pi}(s)$ defined for the appropriate reward setting.

- The above theorem shows the relationship between the policy gradient and the exact critic function.

Policy Gradient With Function Approximation

- The following theorem shows that the state-action value function can be approximated with a certain function, without affecting the unbiasedness of the policy gradient estimate

Theorem

(Policy Gradient with Function Approximation): If the following two conditions are satisfied:

- ① *Function approximator h_w is compatible to the policy*

$$\nabla_w h_w(s, a) = \nabla_\theta \log \pi_\theta(s, a),$$

- ② *Function approximator h_w minimizes the following mean-squared error*

$$\varepsilon = \int_S d^\pi(s) \int_{\mathcal{A}} \pi_\theta(s, a) \{ (Q^\pi(s, a) - h_w(s, a))^2 \},$$

where $\pi_\theta(s, a)$ denotes the stochastic policy, parameterized by θ , then

$$\nabla_\theta J(\theta) = \int_S d^\pi(s) \int_{\mathcal{A}} \nabla_\theta \pi_\theta(s, a) h_w(s, a) da ds.$$


Reducing Variance Using a Baseline

- The policy gradient theorem generalizes to the case where a state-dependent baseline function is taken into account. This can reduce variance, without changing expectation

$$\begin{aligned}\mathbb{E}_{\pi_{\theta}} \{ \nabla_{\theta} \pi_{\theta}(s, a) b(s) \} &= \int_{\mathcal{S}} d^{\pi}(s) \int_{\mathcal{A}} \nabla_{\theta} \pi_{\theta}(s, a) b(s) da ds \\ &= \int_{\mathcal{S}} d^{\pi}(s) b(s) \nabla_{\theta} \int_{\mathcal{A}} \pi_{\theta}(s, a) da ds = 0\end{aligned}$$

$$\nabla_{\theta} J(\theta) = \int_{\mathcal{S}} d^{\pi}(s) \int_{\mathcal{A}} \nabla_{\theta} \pi_{\theta}(s, a) [h_w(s, a) - b(s)] da ds$$

- A good baseline is the state value function
- The policy gradient can be formulated by both the Q function and the **advantage function**


$$A^{\pi_{\theta}}(s, a) = Q^{\pi_{\theta}}(s, a) - V^{\pi_{\theta}}(s)$$

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{\pi_{\theta}}(s, a)]$$

Standard Actor-Critic Algorithms

- If both conditions in the above theorem are met, then the resulting algorithm is equivalent to the REINFORCE algorithm
- Practical actor-critic algorithms often relax the second condition: use TD learning to update the critic approximator.
- TD(0) actor-critic

$$\begin{aligned}\delta_k &= r_k + \gamma V_{w_k}(s_{k+1}) - V_{w_k}(s_k) \\ w_{k+1} &= w_k + \alpha_{c,k} \delta_k \nabla_w V_{w_k}(s_k) \\ \theta_{k+1} &= \theta_k + \alpha_{a,k} \delta_k \nabla_{\theta} \log \pi_{\theta}(s_k, a_k)\end{aligned}$$

- The TD error is actually an estimate of the advantage function

Practical Actor-Critic Variants

- The **policy gradient** has many equivalent forms

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \, v_t] && \text{REINFORCE} \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \, Q^w(s, a)] && \text{Q Actor-Critic} \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \, A^w(s, a)] && \text{Advantage Actor-Critic} \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \, \delta] && \text{TD Actor-Critic}\end{aligned}$$

Natural Policy Gradient

- The vanilla gradient is sensitive to policy parameterizations
- The natural policy gradient is parameterization independent
- It finds ascent direction that is closest to vanilla gradient, when changing policy by a small, fixed amount

$$\nabla_{\theta}^{\text{nat}} \pi_{\theta}(s, a) = G_{\theta}^{-1} \nabla_{\theta} \pi_{\theta}(s, a)$$



Fisher information
matrix (FIM)

$$G_{\theta} = \mathbb{E}_{\pi_{\theta}} \left[\nabla_{\theta} \log \pi_{\theta}(s, a) \nabla_{\theta} \log \pi_{\theta}(s, a)^T \right]$$

- Natural policy gradient methods converges faster in most practical cases. However, estimating the FIM may induce large computation cost

Natural Actor-Critic

- Using compatible function approximation

$$\nabla_w A_w(s, a) = \nabla_\theta \log \pi_\theta(s, a)$$

- The natural policy gradient is then

$$\begin{aligned}\nabla_\theta J(\theta) &= \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(s, a) A^{\pi_\theta}(s, a)] \\ &= \mathbb{E}_{\pi_\theta} \left[\nabla_\theta \log \pi_\theta(s, a) \nabla_\theta \log \pi_\theta(s, a)^T w \right] \\ &= G_\theta w\end{aligned}$$

$$\nabla_\theta^{\text{nat}} J(\theta) = w$$

Deep Reinforcement Learning

- Use deep neural networks to represent
 - Value function
 - Policy
- Optimize loss function by stochastic gradient descent

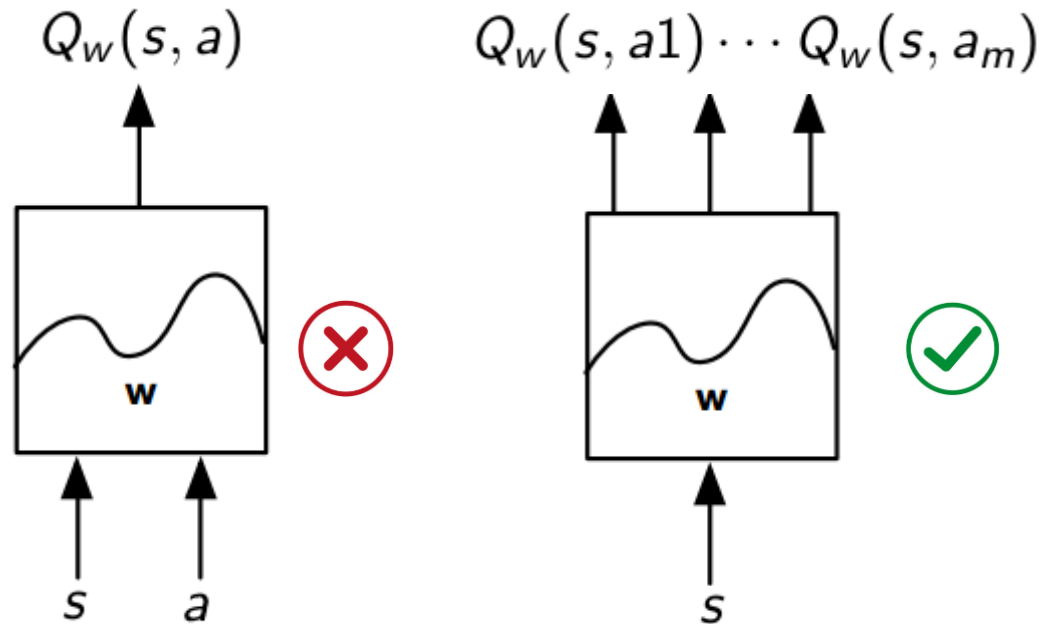
Outline

- Introduction to Deep Learning
- Introduction to Reinforcement Learning (RL)
- **Value-Based Deep RL**
- Policy-Based Deep RL
- Other Deep RL Extensions
- Deep RL Applications

Q-Networks

- Represent the **state-action value function** by **Q-network** with weights w

$$Q_w(s, a) \approx Q^*(s, a)$$



Q-Learning

- Optimal Q-values obey Bellman equation

$$Q^*(s, a) = \mathbb{E}_{s'} \left\{ \underbrace{r + \gamma \max_{a'} Q(s', a')}_{\text{target}} | s, a \right\}$$


- Treat right-hand side as a target and minimize MSE loss by SGD

$$l = \left(r + \gamma \max_{a'} Q_w(s', a') - Q_w(s, a) \right)^2$$

- Converge guarantee using table lookup representation
- But **diverges** using neural networks due to
 - Correlations between samples
 - Non-stationary targets

Deep Q-Networks (DQN)

- Experience replay
 - Build data set from agent's own experience
 - Sample experiences uniformly from data set to remove correlations



- Target Network
 - To deal with non-stationarity, target parameters \hat{w} are held fixed

$$l = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left\{ \left(r + \gamma \max_{a'} Q_{\hat{w}}(s', a') - Q_w(s, a) \right)^2 \right\}$$

Double DQN

- Q-learning is known to overestimate state-action values
 - The max operator uses the same values to select and evaluate an action

$$Q^*(s, a) = \mathbb{E}_{s'} \left\{ r + \gamma \max_{a'} Q(s', a') | s, a \right\}$$

- The upward bias can be removed by decoupling the selection from the evaluation
 - Current Q-network is used to **select** actions
 - Older Q-network is used to **evaluate** actions

$$l = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left\{ \left(r + \gamma Q_{\hat{w}_i}(s', \arg \max_{a'} Q_{w_i}(s', a')) - Q_{w_i}(s, a) \right)^2 \right\}$$

Prioritized Replay

- Uniform experience replay samples transitions regardless of their significance
- Can weight experience according to their significance
- Prioritized replay store experience in priority queue according to the TD error

$$|r + \gamma \max_{a'} Q_{\hat{w}}(s', a') - Q_w(s, a)|$$


- Use **stochastic sampling** to increase sample diversity

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}$$

$$p_i = |\delta_i| + \epsilon$$

Dueling Network

- Dueling network splits Q-network into two channels
 - Action-independent **value function** $V(s)$
 - Action-dependent **advantage function** $A(s, a)$




- The two stream are aggregated to get the Q function

$$Q(s, a; \beta, w_1, w_2) = V(s; \beta, w_1) + \left(A(s, a; \beta, w_2) - \max_{a'} A(s, a'; \beta, w_2) \right)$$

Deep Recurrent Q-Network (DRQN)

- DQNs learn a mapping from a limited number of past states.
- Most real world environments are Partially-Observable Markov Decision Process (POMDP)
- DRQN replaces DQN's first fully connected layer with a LSTM (one variant of RNN)



Asynchronous Q-Learning Variations

- Asynchronous RL
 - Exploits **multithreading** of standard CPU
 - Execute many instances of agent **in parallel**
 - Parallelism **decorrelates** data
 - Thus an alternative to experience replay, which is memory inefficient
 - Network parameters shared between threads
- Asynchronous one-step Q-learning
- Asynchronous one-step SARSA

$$r + \gamma Q(s', a')$$

- Asynchronous n-step Q-learning

$$r_t + \gamma r_{t+1} + \cdots + \gamma^{n-1} r_{t+n-1} + \max_a \gamma^n Q(s_{t+n}, a)$$

Outline

- Introduction to Deep Learning
- Introduction to Reinforcement Learning (RL)
- Value-Based Deep RL
- **Policy-Based Deep RL**
- Other Deep RL Extensions
- Deep RL Applications

Asynchronous Advantage Actor-Critic (A3C)

- Estimate state value function by neural networks

$$V_w(s) \approx \mathbb{E} \{ r_{t+1} + \gamma r_{t+2} + \dots | s \}$$

- Q-value estimated by an n-step sample

$$q_t = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n V_w(s_{t+n})$$

- Actor is updated by advantage policy gradient

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \pi_{\theta}(s_t, a_t) (q_t - V_w(s_t))$$

- Critic is updated by TD learning

$$l_v = (q_t - V_w(s_t))^2$$

Trust Region Policy Optimization (TRPO)

- Formulated as a trust region optimization problem, where each update of the policy is guaranteed to improve

$$\begin{array}{ll} \max_{\pi_{\theta_{\text{old}}}} L_{\pi_{\theta_{\text{old}}}(\pi_{\theta})} & \longrightarrow L_{\pi}(\tilde{\pi}) = J(\pi) + \int_{\mathcal{S}} d^{\pi}(s) \int_{\mathcal{A}} \tilde{\pi}(s, a) A^{\pi}(s, a) \\ \text{subject to } \bar{D}_{KL}(\pi_{\theta_{\text{old}}} || \pi_{\theta}) \leq \delta \end{array}$$

- This provides a unifying perspective on a number of policy update schemes: standard policy gradient, natural policy gradient

$$\max_{\theta} \left[\nabla_{\theta} L_{\pi_{\theta_{\text{old}}}(\pi_{\theta})} |_{\theta=\theta_{\text{old}}} (\theta - \theta_{\text{old}}) \right] \quad \leftarrow \text{First order approximation to the objective}$$

$$\text{subject to } \frac{1}{2}(\theta_{\text{old}} - \theta)^T \mathbf{F}(\theta_{\text{old}})(\theta_{\text{old}} - \theta) \leq \delta$$

$$\mathbf{F}(\theta_{\text{old}})_{ij} = \frac{\partial}{\partial \theta_i} \frac{\partial}{\partial \theta_j} \mathbb{E} \{ D_{KL}(\pi_{\theta_{\text{old}}}(s, \cdot) || \pi_{\theta}(s, \cdot)) \}$$

Natural policy gradient

$$\text{subject to } \frac{1}{2} ||\theta - \theta_{\text{old}}||^2 \leq \delta$$

Standard policy gradient

Practical TRPO Algorithm

- Use the same approximation schemes as the natural policy gradient
- TRPO enforces the constraint by line search
 - Increases stability in practice
- Use a conjugate gradient algorithm to compute the natural gradient direction
 - Makes it practical for deep neural network policies

Deep Deterministic Policy Gradient (DDPG)

- Deterministic policy gradient

$$\nabla_{\theta} J(\pi_{\theta}) = \int_{\mathcal{S}} d^{\pi}(s) \nabla_a Q^{\pi}(s, a)|_{a=\pi_{\theta}(s)} \nabla_{\theta} \pi_{\theta}(s) ds$$

- DDPG is the continuous analogue of DQN

- Experience replay
- Critic estimates value of current policy as in DQN

$$l_w = (r + \gamma Q_{\hat{w}}(s', \pi_{\hat{\theta}}(s')) - Q_w(s, a))^2$$

- Actor updates policy in the deterministic policy gradient direction

$$\nabla_a Q_w(s, a)|_{s=s_t, a=\pi_{\theta}(s_t)} \nabla_{\theta} \pi_{\theta}(s)|_{s=s_t}$$

- The critic provides loss function for the actor

Outline

- Introduction to Deep Learning
- Introduction to Reinforcement Learning (RL)
- Value-Based Deep RL
- Policy-Based Deep RL
- **Other Deep RL Extensions**
- Deep RL Applications

Continuous Q-Learning

- General Q function parameterizations are difficult to find the maximum
- Specific Q function parameterizations can have analytic solution on the maximum

$$A(s, a; w^A) = -\frac{1}{2}(a - \mu(s; w^\mu))^T \mathbf{P}(s; w^P)(a - \mu(s; w^\mu))$$

$$Q(s, a; w^A, w^V) = A(s, a; w^A) + V(s; w^V)$$


- The action that maximizes the Q function is always $\mu(s; w^\mu)$
- The parameterization can then be trained by DQN

Q-Learning with Model-Based Acceleration

- Use model-based methods to generate exploration behavior for Q-learning
 - In practice it often brings very small or no improvement
 - Off-policy model-based exploration is too different from the Q learning policy
- Imagination rollouts
 - Generate synthetic experiences under a learned model by model-based methods
 - Adding synthetic samples to replay buffer
 - Increases sample efficiency in practice

Guided Policy Search (GPS)

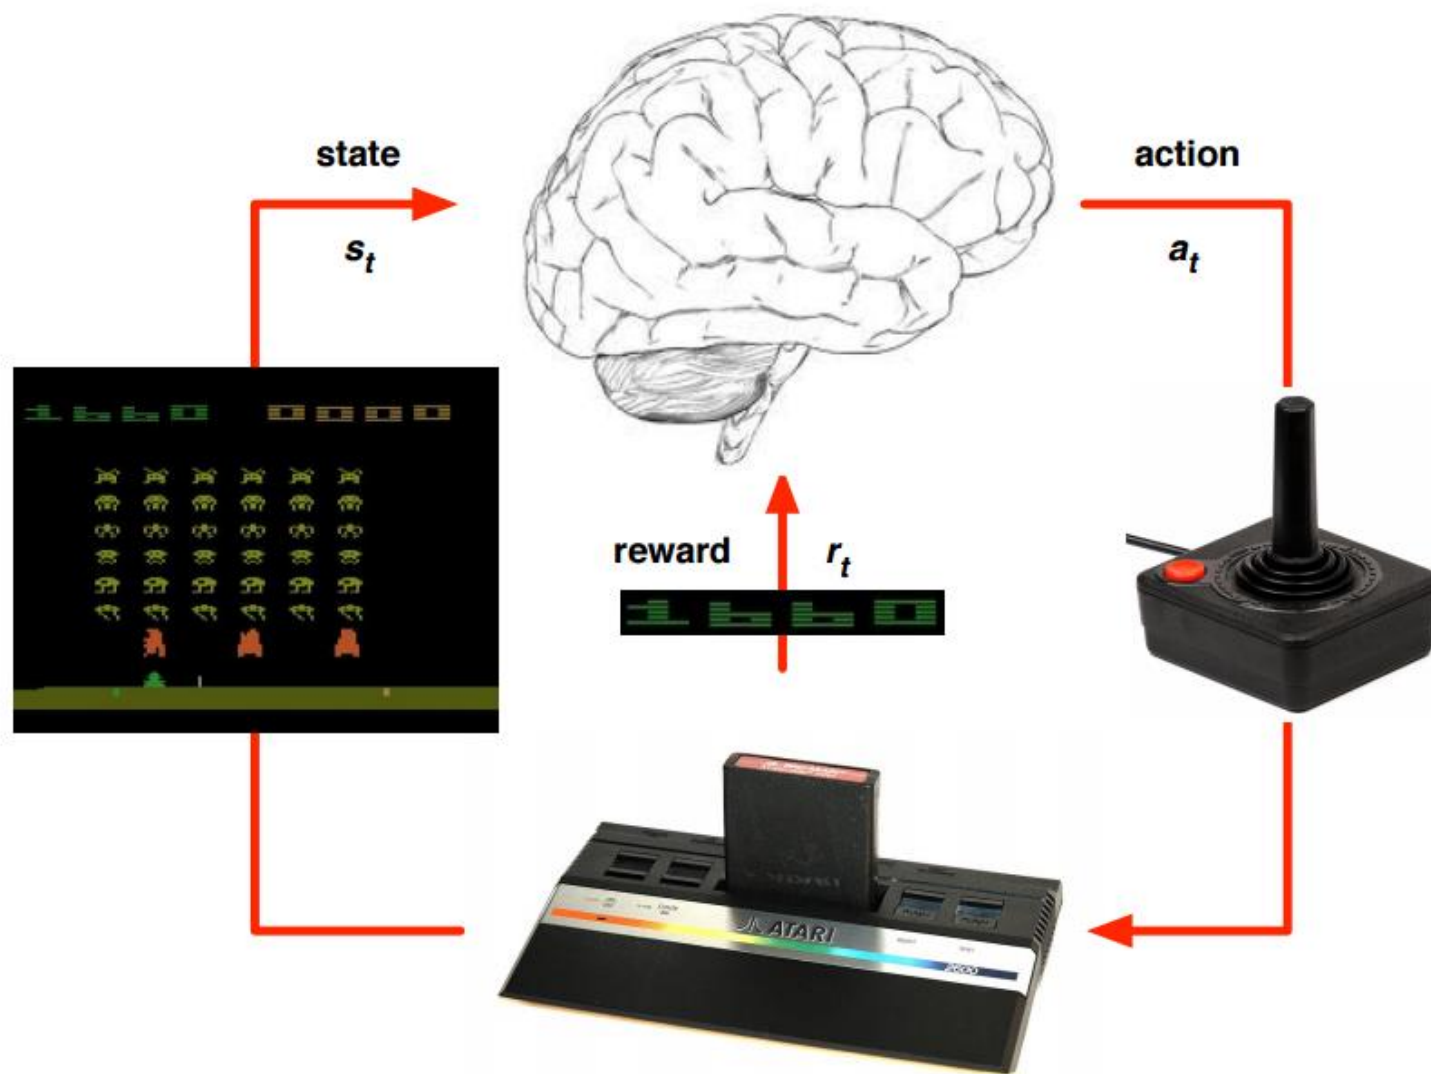
- GPS converts policy search into supervised learning
- Basically a model-based trajectory optimization algorithm generates training data to supervised train the neural network policy
- To enforce convergence, GPS alternates between trajectory optimization and supervised learning
- GPS is data efficient



Outline


- Introduction to Deep Learning
- Introduction to Reinforcement Learning (RL)
- Value-Based Deep RL
- Policy-Based Deep RL
- Other Deep RL Extensions
- **Deep RL Applications**

Deep RL in Atari




DQN in Atari

- End-to-end learning of state-action values from raw pixels
- Input state is stack of raw pixels from last 4 frames
- Output is state-action values from all possible actions
- Reward is change in score for that step




DQN Results in Atari



DQN Variants in Minecraft

- Challenges in environments
 - Partial observability (first-person visual observations)
 - Delayed rewards
 - High-dimensional perception



- Combine DQN with memory network to solve this kind of task

Other Games

- First-person shooter games
- Car racing games




Deep RL in Go


- Use supervised learning followed by deep RL
- Falls in the category of actor-critic framework
 - Use policy network to select moves
 - Use value network to evaluate board positions
- The learning approach is further combined with Monte Carlo search
- It beats the human world champion




Deep RL for Classic Control Tasks



Simple tasks






Locomotion tasks



Hierarchical tasks

Suitable for benchmark various algorithms

Deep RL in Real-World Robotics



Possible Future Research Directions

- Model-free deep RL methods are data inefficient
 - Successful applications focus on game playing and simulator
- How to design more efficient deep RL method
 - Combining with model-based methods or supervised learning
- How to transfer policies trained in simulator to real-world environments
- Transfer learning and multi-task learning of similar policies

